

WIA2005: Algorithm Design and Analysis

Semester 2, Session 2016/17

Lecture 6: Heaps and Heap Sort

Learning Objectives

- Heaps
 - Max-heapify
 - Build-Max-Heap
 - Heapsort
 - Priority queue operation

Heapsort


- Like merge sort, but unlike insertion sort, heapsort's running time is $O(n \lg n)$.
- Like insertion sort, but unlike merge sort, heapsort sorts in place: only a constant number of array elements are stored outside the input array at any time.
- Thus, heapsort combines the better attributes of the two sorting algorithms we have already discussed.
- Heapsort also introduces another algorithm design technique: using a data structure, in this case one we call a “heap,” to manage information.

Heap


- The *(binary) heap* data structure is an array object that we can view as a nearly complete binary tree.
- Each node of the tree corresponds to an element of the array.
- The tree is completely filled on all levels except possibly the lowest, which is filled from the left up to a point.
- An array A that represents a heap is an object with two attributes:
 - *A.length*, which (as usual) gives the number of elements in the array.
 - *A.heap-size*, which represents how many elements in the heap are stored within array A.

Max-heap

- A max-heap viewed as **(a)** a binary tree and **(b)** an array. The number within the circle at each node in the tree is the value stored at that node. The number above a node is the corresponding index in the array. Above and below the array are lines showing parent-child relationships; parents are always to the left of their children. The tree has height three; the node at index 4 (with value 8) has height one.



(a)



(b)

Computing indices

- The root of the tree is $A[1]$, and given the index i of a node, we can easily compute the indices of its parent, left child, and right child:

PARENT(i)

1 **return** $\lfloor i/2 \rfloor$

LEFT(i)

1 **return** $2i$

RIGHT(i)

1 **return** $2i + 1$

- The **LEFT** procedure can compute $2i$ in one instruction by simply shifting the binary representation of i left by one bit position.
- Similarly, the **RIGHT** procedure can quickly compute $2i+1$ by shifting the binary representation of i left by one bit position and then adding in a 1 as the low-order bit.
- The **PARENT** procedure can compute $\lfloor i/2 \rfloor$ by shifting i right one bit position.

Heaps property

- There are two kinds of binary heaps:
 - max-heaps
 - min-heaps

- Max-heaps property:

$$A[\text{PARENT}(i)] \geq A[i]$$

- Min-heaps property:

$$A[\text{PARENT}(i)] \leq A[i]$$

Basic operation on heaps

- MAX-HEAPIFY
- BUILD-MAX-HEAP
- HEAPSORT
- MAX-HEAP-INSERT, HEAP-EXTRACT-MAX, HEAP-INCREASE-KEY, and HEAP-MAXIMUM – for priority queue




Maintaining the heap property

- In order to maintain the max-heap property, we call the procedure MAX-HEAPIFY.

```
MAX-HEAPIFY( $A, i$ )
1   $l = \text{LEFT}(i)$ 
2   $r = \text{RIGHT}(i)$ 
3  if  $l \leq A.\text{heap-size}$  and  $A[l] > A[i]$ 
4       $\text{largest} = l$ 
5  else  $\text{largest} = i$ 
6  if  $r \leq A.\text{heap-size}$  and  $A[r] > A[\text{largest}]$ 
7       $\text{largest} = r$ 
8  if  $\text{largest} \neq i$ 
9      exchange  $A[i]$  with  $A[\text{largest}]$ 
10     MAX-HEAPIFY( $A, \text{largest}$ )
```

Visualizing MAX-HEAPIFY

- The action of $\text{MAX-HEAPIFY}(A, 2)$, where $A.\text{heap-size} = 10$. **(a)** The initial configuration, with $A[2]$ at node $i = 2$ violating the max-heap property since it is not larger than both children.
- The max-heap property is restored for node 2 in **(b)** by exchanging $A[2]$ with $A[4]$, which destroys the max-heap property for node 4.
- The recursive call $\text{MAX-HEAPIFY}(A, 4)$ now has $i = 4$. After swapping $A[4]$ with $A[9]$, as shown in **(c)**, node 4 is fixed up, and the recursive call $\text{MAX-HEAPIFY}(A, 9)$ yields no further change to the data structure.



Running time for MAX-HEAPIFY

- The running time of MAX-HEAPIFY on a subtree of size n rooted at a given node i is the $\Theta(1)$ time to fix up the relationships among the elements $A[i]$, $A[\text{LEFT}(i)]$, and $A[\text{RIGHT}(i)]$, plus the time to run MAX-HEAPIFY on a subtree rooted at one of the children of node i (assuming that the recursive call occurs).
- The running time of MAX-HEAPIFY by the recurrence:

$$T(n) \leq T(2n/3) + \Theta(1)$$

$$T(n) = O(\lg n)$$

Case 2 of Masters Theorem

- Alternatively, we can characterize the running time of MAX-HEAPIFY on a node of height h as $O(h)$.

Building a heap


- We can use the procedure MAX-HEAPIFY in a bottom-up manner to convert an array $A[1, \dots, n]$, where $n = A.length$, into a max-heap.

BUILD-MAX-HEAP(A)

```
1   $A.heap-size = A.length$   
2  for  $i = \lfloor A.length/2 \rfloor$  downto 1  
3      MAX-HEAPIFY( $A, i$ )
```

Visualizing BUILD-MAX-HEAP

The operation of BUILD-MAX-HEAP, showing the data structure before the call to MAX-HEAPIFY in line 3 of BUILD-MAX-HEAP. **(a)** A 10-element input array A and the binary tree it represents. The figure shows that the loop index i refers to node 5 before the call MAX-HEAPIFY(A,i). **(b)** The data structure that results. The loop index i for the next iteration refers to node 4. **(c)–(e)** Subsequent iterations of the **for** loop in BUILD-MAX-HEAP. Observe that whenever MAX-HEAPIFY is called on a node, the two subtrees of that node are both max-heaps. **(f)** The max-heap after BUILD-MAX-HEAP finishes.



Running time for BUILD-MAX-HEAP

- The time required by MAX-HEAPIFY when called on a node of height h is $O(h)$, and so we can express the total cost of BUILD-MAX-HEAP as being bounded from above by

$$\sum_{h=0}^{\lfloor \lg n \rfloor} \left\lceil \frac{n}{2^{h+1}} \right\rceil O(h) = O \left(n \sum_{h=0}^{\lfloor \lg n \rfloor} \frac{h}{2^h} \right)$$

- We evaluate the last summation by substituting $x = 1/2$ in the formula (**Integrating and differentiating series**) yielding

$$\begin{aligned} \sum_{h=0}^{\infty} \frac{h}{2^h} &= \frac{1/2}{(1 - 1/2)^2} \\ &= 2. \end{aligned}$$

- Thus, we can bound the running time of BUILD-MAX-HEAP as

$$\begin{aligned} O \left(n \sum_{h=0}^{\lfloor \lg n \rfloor} \frac{h}{2^h} \right) &= O \left(n \sum_{h=0}^{\infty} \frac{h}{2^h} \right) \\ &= O(n). \end{aligned}$$

Heapsort Algorithm












- The heapsort algorithm starts by using BUILD-MAX-HEAP to build a max-heap on the input array $A[1..n]$, where $n = A.length$.
- Since the maximum element of the array is stored at the root $A[1]$, we can put it into its correct final position by exchanging it with $A[n]$.

HEAPSORT(A)

```
1  BUILD-MAX-HEAP( $A$ )
2  for  $i = A.length$  downto 2
3      exchange  $A[1]$  with  $A[i]$ 
4       $A.heap-size = A.heap-size - 1$ 
5      MAX-HEAPIFY( $A, 1$ )
```

Visualizing HEAPSORT

The operation of HEAPSORT. **(a)** The max-heap data structure just after BUILD-MAX-HEAP has built it in line 1. **(b)–(j)** The max-heap just after each call of MAX-HEAPIFY in line 5, showing the value of i at that time. Only lightly shaded nodes remain in the heap. **(k)** The resulting sorted array A.



Running time for Heapsort

- The HEAPSORT procedure takes time $O(n \lg n)$, since the call to BUILD-MAX-HEAP takes time $O(n)$ and each of the $n - 1$ calls to MAX-HEAPIFY takes time $O(\lg n)$.

Priority Queue

- Heapsort is an excellent algorithm, but a good implementation of quicksort, usually beats it in practice.
- Nevertheless, the heap has its speciality:
 - most popular applications of a heap: as an efficient priority queue.
- Priority queues come in two forms:
 - max-priority queues.
 - min-priority queues.
- A ***priority queue*** is a data structure for maintaining a set S of elements, each with an associated value called a ***key***.

Max-priority and min-priority Queue

- A ***max-priority queue*** supports the following operations:

INSERT(S, x) inserts the element x into the set S , which is equivalent to the operation $S = S \cup \{x\}$.

MAXIMUM(S) returns the element of S with the largest key.

EXTRACT-MAX(S) removes and returns the element of S with the largest key.

INCREASE-KEY(S, x, k) increases the value of element x 's key to the new value k , which is assumed to be at least as large as x 's current key value.

- Alternatively, a ***min-priority queue*** supports the operations INSERT, MINIMUM, EXTRACT-MIN, and DECREASE-KEY.

Max-priority queue – HEAP-MAXIMUM

- The procedure HEAP-MAXIMUM implements the MAXIMUM operation in $O(1)$ time.

```
HEAP-MAXIMUM( $A$ )  
1  return  $A[1]$ 
```

Running time of HEAP-MAXIMUM

- The procedure HEAP-MAXIMUM implements the MAXIMUM operation in $O(1)$ time.

Max-priority queue - HEAP-EXTRACT-MAX

- The procedure HEAP-EXTRACT-MAX implements the EXTRACT-MAX operation.
- It is similar to the **for** loop body (lines 3–5) of the HEAPSORT procedure.

```
HEAP-EXTRACT-MAX(A)  
1  if A.heap-size < 1  
2      error “heap underflow”  
3  max = A[1]  
4  A[1] = A[A.heap-size]  
5  A.heap-size = A.heap-size – 1  
6  MAX-HEAPIFY(A, 1)  
7  return max
```

Running time of HEAP-EXTRACT-MAX

- The running time of HEAP-EXTRACT-MAX is $O(\lg n)$, since it performs only a constant amount of work on top of the $O(\lg n)$ time for MAX-HEAPIFY.

Max-priority queue - HEAP-INCREASE-KEY





- The procedure HEAP-INCREASE-KEY implements the INCREASE-KEY operation.

HEAP-INCREASE-KEY(A, i, key)

```
1  if  $key < A[i]$ 
2      error “new key is smaller than current key”
3   $A[i] = key$ 
4  while  $i > 1$  and  $A[\text{PARENT}(i)] < A[i]$ 
5      exchange  $A[i]$  with  $A[\text{PARENT}(i)]$ 
6       $i = \text{PARENT}(i)$ 
```


Visualizing HEAP-INCREASE-KEY

The operation of HEAP-INCREASE-KEY. **(a)** The max-heap of (a) with a node whose index is i heavily shaded. **(b)** This node has its key increased to 15. **(c)** After one iteration of the **while** loop of lines 4–6, the node and its parent have exchanged keys, and the index i moves up to the parent. **(d)** The max-heap after one more iteration of the **while** loop. At this point, $A[\text{PARENT}(i)] \geq A[i]$. The max-heap property now holds and the procedure terminates.



Running time for HEAP-INCREASE-KEY

- The running time of HEAP-INCREASE-KEY on an n -element heap is $O(\lg n)$, since the path traced from the node updated in line 3 to the root has length $O(\lg n)$.

Max-priority queue - MAX-HEAP-INSERT

- The procedure MAX-HEAP-INSERT implements the INSERT operation.

MAX-HEAP-INSERT(A, key)

1 $A.heap-size = A.heap-size + 1$

2 $A[A.heap-size] = -\infty$

3 HEAP-INCREASE-KEY($A, A.heap-size, key$)


Running time for MAX-HEAP-INSERT

- The running time of MAX-HEAP-INSERT on an n -element heap is $O(\lg n)$.

Summary

- In summary, a heap can support any priority-queue operation on a set of size n in $O(\lg n)$ time.

In the next lecture..



Lecture 7: Data Structure I – Hash Table